

0.1 Mixing issues with the GP and resolutions

We currently use an under-relaxed method to propose new intensity functions when we assume the intensity has a Gaussian Process (GP) prior. In the under relaxed method the candidate intensity function \mathbf{x}^* is composed of the current intensity function \mathbf{x} and a draw from the GP prior ν , namely

$$\log \mathbf{x}^* = \sqrt{1 - \omega^2} \log(\mathbf{x}) + \omega \nu, \quad \nu \sim \mathcal{N}(\mathbf{0}, \Sigma), \quad (1)$$

where ω is a step size parameter. However, we found that this method struggles to capture the behaviour of the intensity before the first spike and after the last spike. In particular, in these regions the posterior distribution of the intensity function often contains a peak in intensity at either time 0 or at T —the end of experiment time — or both. For example in Figure 1(A) we show the posterior distribution of the intensity function given a Ca^{2+} spike sequence — shown as black ticks — where

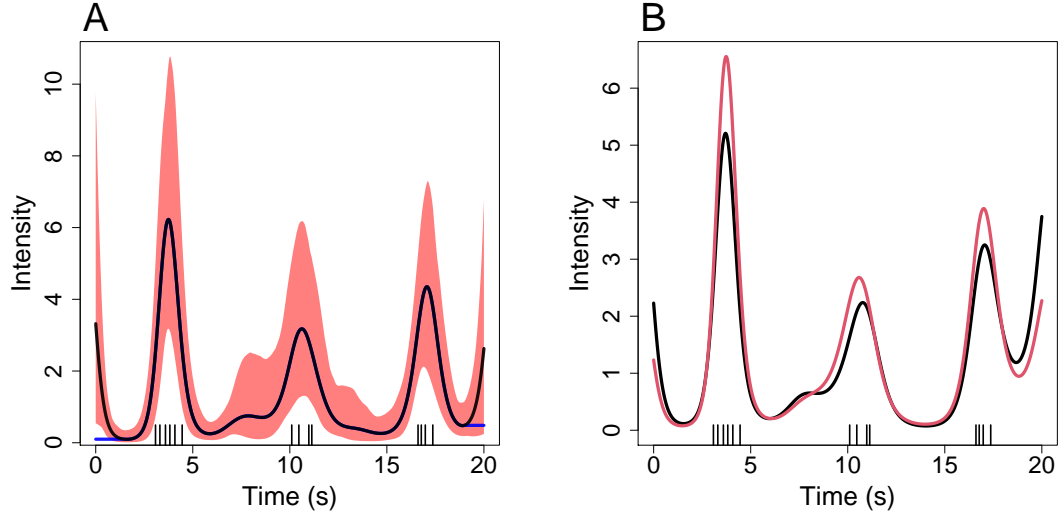


Figure 1: (A) Posterior mean — black line — and 95% confidence interval — red shaded area — for the intensity function. The blue line shows an altered version of the posterior mean which is flat in $[0, 1.5]$ and $[19, 20]$ (B) Two consecutive iterations from the MCMC calculation. The black ticks show the spike times used to infer the intensity function.

we assume a GP prior and a Gamma ISI distribution. The inference was calculated using 80000 iterations after the initial 50000 were binned, where time was discretised into 2000 steps, the GP hyper parameters were set to $\sigma_f^2 = 1000$, $\sigma_n^2 = 1e^{-5}$, $l = 1.59$, the under-relaxed parameter $\omega = 0.001$ and the prior for the ISI parameter was an exponential with rate 0.01. We see that the posterior intensity distribution captures the three peaks in intensity where the Ca^{2+} spikes occur. Furthermore, we see that the intensity does decrease before the first peak and after the last peak in intensity — approximately at the times $[1.6, 3.7]$ and $[17, 19]$. However, we also find large values of intensity at times 0s and 20s. Since there are no spikes close to these times you would expect the intensity to be close to zero. To quantify if our posterior distribution should have these regions of high intensity we compare the posterior mean with an altered version where the intensity is flattened in regions $[0, 1.5]$ and $[19, 20]$ — shown in blue in Figure 1. We found that the log likelihood of the posterior mean is -8.90 and the altered version is -7.20 . Therefore, we find that there is higher probability of these regions being flat.

The problem in posterior computation lies in the mixing of the under-relaxed method. Since the candidate intensity varies over all time the change in one region may outweigh a change in another region. In particular, in Figure 1(B) we give an example of an iteration where the current intensity function \mathbf{x} is shown by the red line and the candidate intensity function \mathbf{x}^* by the black line. We see that in $[0, 1]$ and $[19, 20]$ \mathbf{x}^* is much greater than \mathbf{x} . However, \mathbf{x}^* is smaller than \mathbf{x} in the three peaks — at times 4s, 10s and 17s — which is closer to the underlying rate. Therefore we find the log likelihood of \mathbf{x} is -8.13 and \mathbf{x}^* is -7.78 . Therefore, we would accept the candidate function. Moreover, once the features in $[0, 1]$ and $[19, 20]$ are accepted they are difficult to remove since a candidate function is likely to perform poorly in other regions of the function. Therefore, we want to create a new proposal mechanism to work in tandem with the under-relaxed method to improve the mixing at the beginning and end of the intensity function. We will do this by proposing candidate intensity functions that vary only on these regions of interest. In particular, we will call the proposals the conditional method at the start and conditional method at the

end, for the proposals that only change values in the beginning of the experiment
and end of experiment respectively.

0.1.1 Method

We begin by updating the intensity function only at the beginning of the function.
Assume initially that we have discretised time into $N + 1$ steps by $\mathbf{t} = \{t_0, t_1, \dots, t_N\}$
defined as $t_i = iT/N$ for $i \in \{0, \dots, N\}$. Furthermore, the current value of the
discretised intensity function is $\mathbf{x} = \{x_0, \dots, x_N\}$ where $x_i = x(t_i)$ for $i \in \{0, \dots, N\}$.
We need to decide how to partition \mathbf{x} into two regions — the part we update and the
part that remains unchanged. We choose to sample the partition value uniformly
between time zero and just after the first spike time. We do this because this region is
where the under-relaxed method has difficulties mixing. Suppose the first spike occurs
at time t_ι , where $\iota \in \{1, \dots, N\}$. Then the partition value M is drawn uniformly
from the set $\{1, 2, \dots, \min(\iota + w, N)\}$, where w controls how far past the first spike
we can sample. The minimum is used to only allow valid index values. We split time
into two groups; one either side of t_M — $A = \{t_0, \dots, t_M\}$ and $B = \{t_{M+1}, \dots, t_N\}$. We
similarly split the \mathbf{x} into two groups $\mathbf{x}_A = \{x_0, \dots, x_M\}$ and $\mathbf{x}_B = \{x_{M+1}, \dots, x_N\}$.
Therefore, the candidate intensity function \mathbf{x}^* consists of two parts: \mathbf{x}_A^* describing
the new values we propose in region A and \mathbf{x}_B^* which corresponds to the unchanged
values in region B , $\mathbf{x}_B^* = \mathbf{x}_B$.

In Figure 2 we give an example of splitting \mathbf{x} — the black line — into two regions.
In this example \mathbf{x} is discretised into 2000 steps. The first spike occurs at 3.08s which
corresponds to the 308th element of \mathbf{t} . Taking $w = 100$ we sample the partition index
from $\{1, \dots, 408\}$ which corresponds to times $\{0.01, \dots, 4.08\}$ — which is shown as
the grey box. A realisation is shown by the red dotted line where $M = 326$. Thus,
the thin black line in region B will remain the same for the candidate function and
we propose new values in region A .

We now need to choose a method to propose \mathbf{x}_A^* given the current value \mathbf{x}_A . A natural
option would be to propose candidate values from a multivariate normal (MVN)
distribution. Recall that we say the random vector $\mathbf{X} = \{X_0, X_1, \dots, X_M\}$ comes




Figure 2: Illustration of how to partition \mathbf{x} into two regions A and B . The partition value is drawn from the grey region. An example draw t_M is shown by the dotted red line. The thick black line — left of t_M — is the region of \mathbf{x} we will propose new values for, whereas the thin black line — right of t_M — will remain unchanged. The black ticks on the x-axis represent spike times.

from a $(M + 1)$ -dimensional MVN distribution $\mathbf{X} \sim \mathcal{N}_{M+1}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ if its probability density function $\mathcal{N}_{M+1}(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma})$ of $\mathbf{x} = \{x_0, \dots, x_M\}$ is given by

$$\mathcal{N}_{M+1}(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma}) = (2\pi)^{(M+1)/2} \det(\boldsymbol{\Sigma})^{-1/2} \exp \left[-\frac{1}{2}(\mathbf{X} - \boldsymbol{\mu})\boldsymbol{\Sigma}^{-1}(\mathbf{X} - \boldsymbol{\mu}) \right],$$

where $\boldsymbol{\mu}$ is the vector of means and $\boldsymbol{\Sigma}$ is the covariance matrix. In our case we want our proposal to depend on the current value of the intensity function. We do this by making the mean of the MVN depend on \mathbf{x}_A via the function f giving $\boldsymbol{\mu} = f(\mathbf{x}_A)$. The choice of f is discussed in Section 0.1.2. Recall that the prior for the intensity function is a GP with square exponential covariance with hyper parameters $\{\sigma_f^2, \sigma_n^2, l\}$ representing the signal variance, noise, and length scale, respectively. We want our

proposal to match the ‘shape’ of the GP prior, therefore the covariance matrix Σ for our proposal inherits the length scale and noise parameter of the GP. Formally, $\Sigma = \{\Sigma_{i,j}\}_{i,j \in \{1,M\}}$ where

$$\Sigma_{i,j} = \sigma^2 \exp \left[\frac{(t_i - t_j)^2}{2l^2} \right] + \delta(t_i - t_j) \sigma_n^2,$$

68 where l and σ_n^2 come from the GP prior. The variance σ^2 is often taken small, such
69 that the proposed values are close to \mathbf{x}_A .

70 Recall that in the under-relaxed proposal mechanism the proposals are drawn on
71 the logarithmic scale. This is because the intensity function is defined as non-negative.
72 Clearly, the intensity function restricted to region A must also be non-negative.
73 Therefore our proposal also needs to be calculated on the logarithmic scale $\mathbf{x}_A^* = e^{\mathbf{x}_A^\dagger}$
74 where $\mathbf{x}_A^\dagger \sim \mathcal{N}_{M+1}(f(\mathbf{x}_A), \Sigma)$ where the logarithm is subsumed into f .

However, if we choose to propose new values in this manner there is no guarantee that the new values agree with the intensity in region B . In other words, the proposed function would be discontinuous at time t_M . To mitigate this issue we will propose candidate values in the region A by drawing from the MVN distribution described above conditioned on the next K values coming from \mathbf{x}_B . The conditioned MVN distribution follows a MVN distribution with mean and covariance matrix outlined below []. To construct the conditional MVN we first need a MVN distribution defined on region A and the values we wish to condition on, namely $C = \{t_{M+1}, \dots, t_{M+K}\}$ and $\mathbf{x}_C = \{x_{M+1}, \dots, x_{M+K}\}$. It follows that $\mathbf{x}_{A \cup C}^\dagger \sim \mathcal{N}_{M+K+1}(f(\mathbf{x}_{A \cup C}), \Sigma)$. By partitioning into A and C the MVN distribution of $\mathbf{x}_{A \cup C}^\dagger$ can be expressed as

$$\mathbf{x}_{A \cup C}^\dagger = \begin{bmatrix} \mathbf{x}_A^\dagger \\ \mathbf{x}_C^\dagger \end{bmatrix} \sim \mathcal{N}_{M+K+1} \left(\begin{bmatrix} f(\mathbf{x}_A) \\ f(\mathbf{x}_C) \end{bmatrix}, \begin{bmatrix} \Sigma_{11} & \Sigma_{12} \\ \Sigma_{21} & \Sigma_{22} \end{bmatrix} \right),$$

75 where $f(\mathbf{x}_{A \cup C})$ is partitioned into $f(\mathbf{x}_A)$ and $f(\mathbf{x}_C)$. Σ has been partitioned into four
76 blocks: Σ_{11} of size $(M+1) \times (M+1)$, Σ_{12} of size $(M+1) \times K$, Σ_{21} of size $(M+1) \times K$,
77 and Σ_{22} of size $K \times K$. Then the conditional distribution of $\mathbf{x}_{A \cup C}^\dagger | \mathbf{x}_C^\dagger = \log \mathbf{x}_C$ is a
78 $(M+1)$ -dimensional MVN distribution with mean $\bar{\mu}$ and covariance $\bar{\Sigma}$ given by

$$\bar{\mu} = f(\mathbf{x}_A) + \Sigma_{12} \Sigma_{22}^{-1} (\log \mathbf{x}_C - f(\mathbf{x}_C)) \quad (2)$$

and

$$\bar{\Sigma} = \Sigma_{11} - \Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21}.$$

Therefore, we propose candidate values \mathbf{x}_A^* on region A by $\mathbf{x}_A^* = e^{\mathbf{x}_A^\dagger}$ where $\mathbf{x}_A^\dagger \sim \mathcal{N}_{M+1}(\bar{\boldsymbol{\mu}}, \bar{\Sigma})$. Thus the candidate function $\mathbf{x}^* = \mathbf{x}_A^* \frown \mathbf{x}_B$, where $\mathbf{a} \frown \mathbf{b}$ is defined as the concatenation of vectors \mathbf{a} and \mathbf{b} . Simply put, the candidate function consists of a draw from a MVN distribution in region A conditioned on the first K values in B and is equal to the current intensity on all other points — those in B .

The probability that we accept the candidate function is given by

$$p_{\text{acc}} = \frac{\text{Likelihood}(\mathbf{x}^*)\text{Prior}(\mathbf{x}^*)q(\mathbf{x}^*|\mathbf{x})}{\text{Likelihood}(\mathbf{x})\text{Prior}(\mathbf{x})q(\mathbf{x}|\mathbf{x}^*)}.$$

Although often the acceptance probability is calculated on the logarithmic scale. The likelihood is defined in ?? . We have a Gaussian Process prior — with mean zero and square exponential covariance function with hyper parameters $\sigma_f^2, \sigma_n^2, l$ — for the logarithm of our intensity function. Therefore on our discretised time we have $\text{Prior}(\mathbf{x}) = \mathcal{N}_{N+1}(\log \mathbf{x}; \mathbf{0}, \Sigma_{\text{prior}})$, where Σ_{prior} denoted the covariance matrix on \mathbf{t} for the GP. Since the hyper-parameters remain the same for both \mathbf{x}^* and \mathbf{x} the prior ratio simplifies

$$\begin{aligned} \frac{\text{Prior}(\mathbf{x}^*)}{\text{Prior}(\mathbf{x})} &= \frac{(2\pi)^{(N+1)/2} \det(\Sigma_{\text{prior}})^{-1/2} \exp\left[-\frac{1}{2}(\log \mathbf{x}^*)\Sigma_{\text{prior}}^{-1}(\log \mathbf{x}^*)\right]}{(2\pi)^{(N+1)/2} \det(\Sigma_{\text{prior}})^{-1/2} \exp\left[-\frac{1}{2}(\log \mathbf{x})\Sigma_{\text{prior}}^{-1}(\log \mathbf{x})\right]} \\ &= \frac{\exp\left[-\frac{1}{2}(\log \mathbf{x}^*)\Sigma_{\text{prior}}^{-1}(\log \mathbf{x}^*)\right]}{\exp\left[-\frac{1}{2}(\log \mathbf{x})\Sigma_{\text{prior}}^{-1}(\log \mathbf{x})\right]}. \end{aligned}$$

Working on the logarithmic scale gives

$$\log\left(\frac{\text{Prior}(\mathbf{x}^*)}{\text{Prior}(\mathbf{x})}\right) = \frac{1}{2} \left((\log \mathbf{x})\Sigma_{\text{prior}}^{-1}(\log \mathbf{x}) - (\log \mathbf{x}^*)\Sigma_{\text{prior}}^{-1}(\log \mathbf{x}^*) \right).$$

$q(\mathbf{x}^*|\mathbf{x})$ is the proposal distribution — the conditional distribution of proposing function \mathbf{x}^* given \mathbf{x} . For our proposal we have $q(\mathbf{x}^*|\mathbf{x}) = \mathcal{N}_{M+1}(\log \mathbf{x}_A^*; \bar{\boldsymbol{\mu}}(\mathbf{x}_{AUC}), \bar{\Sigma})$, where we explicitly state the dependency of $\bar{\boldsymbol{\mu}}$ on \mathbf{x}_{AUC} . Similar to the prior ratio, we can simplify the proposal ratio since $\bar{\Sigma}$ is the same for $q(\mathbf{x}^*|\mathbf{x})$ and $q(\mathbf{x}|\mathbf{x}^*)$. On

the logarithmic scale this gives

$$\log \left(\frac{q(\mathbf{x}^*|\mathbf{x})}{q(\mathbf{x}|\mathbf{x}^*)} \right) = \frac{1}{2} \left((\log \mathbf{x}_A - \bar{\boldsymbol{\mu}}(\mathbf{x}_{A \cup C}^*)) \bar{\boldsymbol{\Sigma}}^{-1} (\log \mathbf{x}_A - \bar{\boldsymbol{\mu}}(\mathbf{x}_{A \cup C}^*)) \right. \\ \left. - (\log \mathbf{x}_A^* - \bar{\boldsymbol{\mu}}(\mathbf{x}_{A \cup C})) \bar{\boldsymbol{\Sigma}}^{-1} (\log \mathbf{x}_A^* - \bar{\boldsymbol{\mu}}(\mathbf{x}_{A \cup C})) \right). \quad (3)$$

84 To propose intensity functions that vary at the end of the experiment time a
85 symmetric argument can be used where we swap the definitions of regions A and B .

86 0.1.2 Choice of f

In this section we discuss the choice of function f used to proposal \mathbf{x}_A^* . Recall the idea behind these proposals is to ‘flatten’ the intensity in region A because the under-relaxed method alone often struggles to capture this shape. A natural starting point is to use the current intensity giving $f(\mathbf{x}_A) = f_{\text{cur}} = \log \mathbf{x}_A$ — where the log arises from proposing on the logarithmic scale — as the mean of the proposed MVN distribution. This is similar to the classical Metropolis-Hastings algorithm [1]. One advantage of this approach is the acceptance probability simplifies, due to the proposal ratio cancelling out. With this choice of f , the mean of the MVN simplifies due to the second term in (2) cancelling, leaving $\bar{\boldsymbol{\mu}} = \log \mathbf{x}_A$. Substituting into equation (3) we get

$$\log \left(\frac{q(\mathbf{x}^*|\mathbf{x})}{q(\mathbf{x}|\mathbf{x}^*)} \right) = \frac{1}{2} \left((\log \mathbf{x}_A - \log \mathbf{x}_A^*) \bar{\boldsymbol{\Sigma}}^{-1} (\log \mathbf{x}_A - \log \mathbf{x}_A^*) \right. \\ \left. - (\log \mathbf{x}_A^* - \log \mathbf{x}_A) \bar{\boldsymbol{\Sigma}}^{-1} (\log \mathbf{x}_A^* - \log \mathbf{x}_A) \right), \\ = 0,$$

Since $\bar{\boldsymbol{\Sigma}}$ does not depend on \mathbf{x}_A . Thus, the acceptance ratio simplifies to

$$p_{\text{acc}} = \frac{\text{Likelihood}(x_{\text{can}}) \text{Prior}(x_{\text{can}})}{\text{Likelihood}(x_{\text{cur}}) \text{Prior}(x_{\text{cur}})}.$$

87 However, this approach may be slow at ‘flattening’ the intensity in region A . This
88 is because the candidate intensity functions are centered on \mathbf{x}_A . Therefore, we also
89 propose both functions that have ‘steeper’ intensity or ‘flatter’ intensity in region




Figure 3: Illustration of proposal mechanism when $f = \log \mathbf{x}_A$. (A) shows ten candidate functions — grey lines — where the current intensity is shown by the red line. (B) 100 iterations of the proposal mechanism where the original \mathbf{x}_A is the red line. The iterations progress from black to blue. The black ticks show the spike times.

90 A. This can be seen in Figure 3(A) which shows ten candidate intensities — with
 91 fixed partition — zoomed in on the region $[0, 3.3]$. The proposal parameters are
 92 $K = 10$, $M = 337$, $l = 1.59$, $\sigma_n^2 = 1e^{-5}$ and $\sigma^2 = 1$. The red line shows the current
 93 intensity function \mathbf{x}_A and the grey lines show ten candidate intensity functions. We
 94 see that four of the candidate functions have larger intensity than \mathbf{x}_A and the majority
 95 of proposals are close to \mathbf{x}_A . It is also important to note that the smaller \mathbf{x}_A the
 96 more difficult it is to proposal smaller values since the calculations are done on the
 97 logarithmic scale. For example, suppose on the logarithmic scale the intensity at
 98 time t is 0 and the drawn candidate value equals -1 , this translates to an original
 99 intensity of 1 and candidate intensity equal 0.368, hence squashing smaller values
 100 closer together. This implies that it will take many iterations to flatten the curve
 101 in the region $[0, 1]$. Indeed, this is shown in Figure 3(B) where the red line shows
 102 \mathbf{x}_A , we then compute 100 iterations where $K = 10$, $w = 100$, $l = 1.59$, $\sigma_n^2 = 1e^{-5}$
 103 and $\sigma^2 = 1$. The iterations are shown going from the black lines to blue lines as

iterations progress. We see that we do indeed steadily flatten the intensity function in the region $[0, 1]$.

So we have found that using $f = f_{\text{cur}}$ does reduce the peak intensity found at time 0s. However, it often requires a large number of iterations to do this. With this in mind, it may be more intuitive to propose functions whose shape is flatter than the current intensity function. For example proposing intensity functions whose mean is constant. We shall consider two alternative functions for f both of which are constant. Namely, $f_{\text{mean}} = \{\text{mean}(\log \mathbf{x}_A)\}_{i=0}^M$ which returns the constant vector where each element is the mean of $\log \mathbf{x}_A$ and $f_{\text{min}} = \{\min(\log \mathbf{x}_A)\}_{i=0}^M$ which returns the constant vector where each element is the minimum of $\log \mathbf{x}_A$.

In Figures 4(A) and (C) we show ten candidate functions from the proposal distribution using f_{min} and f_{max} , respectively, where the blue lines show the mean of the proposal distribution. This illustrates how using f_{min} and f_{max} makes the candidate intensity functions flatter than those using $f = f_{\text{cur}}$ in Figure 3. For both the proposal parameters were $K = 10$, $M = 337$, $l = 1.59$, $\sigma_v^2 = 1e^{-5}$ and $\sigma^2 = 0.5$. Therefore, if accepted we require fewer iterations than f_{cur} to remove the peak intensity at time 0. We also see a marginal difference between using f_{mean} and f_{min} where the proposed functions using f_{min} are smaller than those using f_{mean} . In (B,D) we plot 100 iterations of using this proposal mechanism, where we begin at the red line. Iterations are shown going from grey to blue lines at the iterations progress. We see that both options flatten the intensity, in a small iterations. For f_{mean} we see a couple iterations remain around the initial intensity before jumping to a flat intensity. This is because the width is small in the initial iterations— for example corresponding to time 0.5 — thus the proposed intensities remain high. However, it took f_{min} and f_{mean} 14 and 15 iterations, respectively, to accept a candidate function with a flat intensity in $[0, 2]$. Therefore, we find that the acceptance probability is lower for f_{mean} and f_{min} compared to f_{cur} . However, it takes f_{cur} far longer to propose functions with initial flat intensity.

In the above discussion we have always used the same initial intensity function, and we found that all considered f do reduce the high intensity found in the region

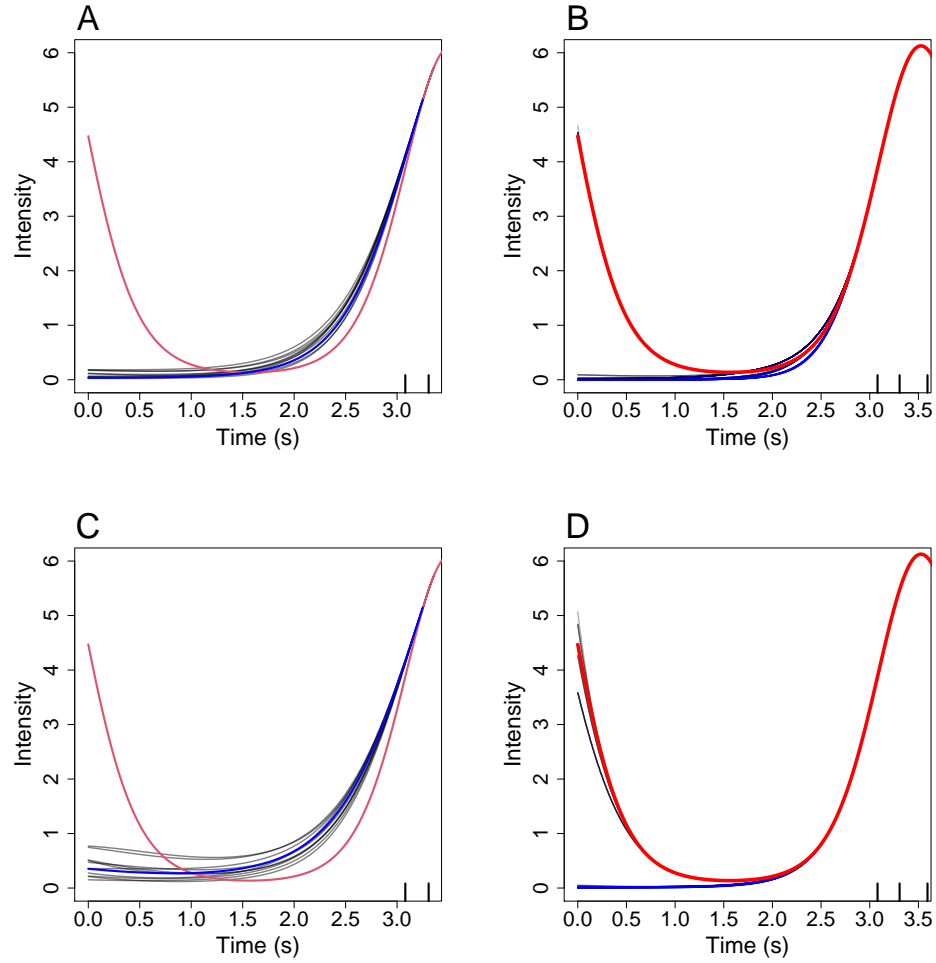


Figure 4: Illustration of the mean of a MVN distribution where black shows the mean for the current version and also the current value of \mathbf{x}_A red uses the mean and blue the min. (A-D) shows the results for different widths of region A where the width gets wider as (A) goes to (D).

134 $[0, 1.5]$. Therefore when deciding which method to implement we are interested in the
 135 speed at which the proposals flatten this region. Therefore, in the particular example
 136 shown one would choose to use either f_{\min} or f_{mean} as they took a similar amount
 137 of iterations. However, in general it can take a large number of iterations to accept
 138 a proposal when using f_{\min} or f_{mean} . For example in Figure 5 we show the mean of
 139 the proposal — the red line — compared to the current intensity function on the log




Figure 5: Example where the mean of the proposal —red line— is far from the current intensity on the logarithmic scale.

140 scale. We see that the two lines differentiate by a large margin, and in this case the
 141 prior ratio stops candidate functions been accepted. In particular, the log prior ratio
 142 of the posterior mean against the current intensity is -54.8 , which outweighs the log
 143 likelihood ratio of 1.08 . The mean of the proposal is decreases quickly because of the
 144 gradient of the values it is conditioned on.

145 Therefore, we find that if candidate functions are accepted fast enough f_{\min} and
 146 f_{mean} will flatten the region quicker. However, there is no guarantee that candidate
 147 functions will get accepted. Thus, although f_{cur} reduces the intensity less per iteration
 148 it will not get stuck not accepting proposals.

149 Other functions could be used in this proposal mechanism, and further research
 150 could look into whether there is an optimal f for proposing updates to the intensity
 151 function on these regions of interest.

Algorithm 1: Inference for GP intensity function.

Input:

Spikes, ISI distribution, T , Numer of iterations N_{iter} , \mathbf{t} , GP hyper parameters, ISI parameter

Output: M samples of intensity \mathbf{x}

Set initial values $\mathbf{x}^{(0)}$;

for i in 1 to $(M + M_{\text{burn}})$ **do**

 Propose \mathbf{x}_{can} and calculate acceptance probability using under-relaxed method, update $\mathbf{x}^{(i)}$;

if $i \% 1000 == 0$ **then**

for j in 1 to 50 **do**

 Use conditional method at the start with f_{min} , and conditional method at the end with f_{min} to update $\mathbf{x}^{(i)}$;

if *No proposals using f_{min} are accepted* **then**

for j in 1 to 50 **do**

 Use conditional method at the start with f_{cur} , and conditional method at the end with f_{cur} to update $\mathbf{x}^{(i)}$;

Return: \mathbf{x} ;

0.1.3 Proposal mechanism for GP

We now need to decide how to use this new proposal in conjunction with the under-relaxed method. We choose to use both f_{cur} and f_{min} because f_{min} can quickly flatten the intensity if accepted, and f_{cur} can reduce the intensity if f_{min} fails to. We only want to update on at the beginning and end of the function infrequently. This is because we are only using this method to help explore regions that the under-relaxed methods struggles to search. Therefore, we will only use this proposal every 1000th iteration, but we shall use the proposal multiple times. This is to improve the chance of removing undesired features in these regions. We choose to initially do 50 iterations on both the beginning and end of the function using f_{min} . If no candidate functions are accepted under f_{min} then we do a further 50 iterations using f_{cur} . The algorithm for the MCMC is shown in Algorithm 1.